

Porras[hw2]

September 3, 2024

Download the dataset from: <https://github.com/bellawillrise/Introduction-to-Numerical-Computing-in-Python/>

Submit a pdf file, which is a rendered saved version of the jupyter notebook. Make sure to execute all the codes so the output can be viewed in the pdf.

Also include the link to the public github repository where the jupyter notebook for the assignment is uploaded.

Link to the github repository: <https://github.com/Megunut/CMSC-197-hw2> .

```
[5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[5]: # %matplotlib inline
```

```
[6]: data = pd.read_csv("data/movie_metadata_cleaned.csv")
```

```
[7]: data.head(2)
```

```
[7]: Unnamed: 0      movie_title  color \
0      0      b'Avatar'  Color
1      1  b"Pirates of the Caribbean: At World's End"  Color

      director_name  num_critic_for_reviews  duration  director_facebook_likes \
0  James Cameron      723.0      178.0      0.0
1  Gore Verbinski      302.0      169.0      563.0

      actor_3_facebook_likes  actor_2_name  actor_1_facebook_likes  ... \
0      855.0  Joel David Moore      1000.0  ...
1      1000.0  Orlando Bloom      40000.0  ...

      num_user_for_reviews  language  country  content_rating  budget \
0      3054.0  English  USA      PG-13  237000000.0
1      1238.0  English  USA      PG-13  300000000.0

      title_year  actor_2_facebook_likes  imdb_score  aspect_ratio \
```

0	2009.0	936.0	7.9	1.78
1	2007.0	5000.0	7.1	2.35

	movie_facebook_likes
0	33000.0
1	0.0

[2 rows x 29 columns]

0.1 Get the top 10 directors with most movies directed and use a boxplot for their gross earnings


```
[9]: # displays a list of the top 10 directors with most movies
top_10_directors = data.groupby("director_name").size().
    ↪sort_values(ascending=False)[1:11]
top_10_directors
```

```
[9]: director_name
Steven Spielberg      26
Woody Allen           22
Clint Eastwood        20
Martin Scorsese       20
Ridley Scott          17
Steven Soderbergh     16
Spike Lee             16
Tim Burton            16
Renny Harlin          15
Oliver Stone          14
dtype: int64
```

```
[11]: top_10_directors = data.groupby("director_name").size().
    ↪sort_values(ascending=False)[1:11]
top_10_directors

# retrieves the data from dataset which match the selected directors
top_10_directors = top_10_directors.index
top_10_directors_data = data[data['director_name'].isin(top_10_directors) &
    ↪data['gross']]
top_10_directors_data

# creates a boxplot of gross earnings of the top 10 directors
plt.figure(figsize=(15,4))
sns.boxplot(x="director_name", y="gross", data=top_10_directors_data)
plt.xticks(rotation=45)
plt.show()
```



0.2 Plot the following variables in one graph:

- num_critic_for_reviews
- IMDB score
- gross

```
[65]: plt.figure(figsize=(15,6))


# create a copy of the dataset, so that future questions won't be affected
# after standardizing
standardized_data = data.copy()

# Standardizing the three variables Z-Score using pandas
standardized_data['num_critic_for_reviews'] = (
    (standardized_data['num_critic_for_reviews'] -
    standardized_data['num_critic_for_reviews'].mean()) /
    standardized_data['num_critic_for_reviews'].std()
standardized_data['imdb_score'] = (standardized_data['imdb_score'] -
    standardized_data['imdb_score'].mean()) / standardized_data['imdb_score'].
    std()
standardized_data['gross'] = (standardized_data['gross'] -
    standardized_data['gross'].mean()) / standardized_data['gross'].std()

# Plots histogram from genres
ax = sns.histplot(standardized_data["num_critic_for_reviews"],bins=30,
    color="red", label="num_critic_for_reviews",linewidth=0, alpha=0.2, )
sns.histplot(standardized_data["imdb_score"],bins=30, color="teal",
    label="imdb_score", linewidth=0, ax=ax, alpha=0.2)
sns.histplot(standardized_data["gross"],bins=30, color="blue", label="gross",
    linewidth=0, ax=ax, alpha=0.2)
```

```
ax.legend()
```

```
[65]: <matplotlib.legend.Legend at 0x1c28275f2c0>
```



0.3 Compute Sales (Gross - Budget), add it as another column

```
[15]: data['sales'] = data['gross'] - data['budget']
data
```

```
[15]:
```

	Unnamed: 0	movie_title	color	\
0	0	b'Avatar'	Color	
1	1	b"Pirates of the Caribbean: At World's End"	Color	
2	2	b'Spectre'	Color	
3	3	b'The Dark Knight Rises'	Color	
4	4	b'Star Wars: Episode VII - The Force Awakens ...	0	
...	
5039	5039	b'The Following	Color	
5040	5040	b'A Plague So Pleasant'	Color	
5041	5041	b'Shanghai Calling'	Color	
5042	5042	b'My Date with Drew'	Color	
5043	5043	b'Starting Over Again'	0	

	director_name	num_critic_for_reviews	duration	\
0	James Cameron	723.0	178.0	
1	Gore Verbinski	302.0	169.0	
2	Sam Mendes	602.0	148.0	
3	Christopher Nolan	813.0	164.0	
4	Doug Walker	0.0	0.0	
...	
5039	0	43.0	43.0	
5040	Benjamin Roberds	13.0	76.0	

5041	Daniel Hsia	14.0	100.0
5042	Jon Gunn	43.0	90.0
5043	Olivia Lamasan	0.0	0.0

	director_facebook_likes	actor_3_facebook_likes	actor_2_name	\
0	0.0	855.0	Joel David Moore	
1	563.0	1000.0	Orlando Bloom	
2	0.0	161.0	Rory Kinnear	
3	22000.0	23000.0	Christian Bale	
4	131.0	0.0	Rob Walker	
...	
5039	0.0	319.0	Valorie Curry	
5040	0.0	0.0	Maxwell Moody	
5041	0.0	489.0	Daniel Henney	
5042	16.0	16.0	Brian Herzlinger	
5043	0.0	0.0	Toni Gonzaga	

	actor_1_facebook_likes	...	language	country	content_rating	\
0	1000.0	...	English	USA	PG-13	
1	40000.0	...	English	USA	PG-13	
2	11000.0	...	English	UK	PG-13	
3	27000.0	...	English	USA	PG-13	
4	131.0	...	0	0	0	
...	
5039	841.0	...	English	USA	TV-14	
5040	0.0	...	English	USA	0	
5041	946.0	...	English	USA	PG-13	
5042	86.0	...	English	USA	PG	
5043	0.0	...	0	Philippines	PG	

	budget	title_year	actor_2_facebook_likes	imdb_score	aspect_ratio	\
0	237000000.0	2009.0	936.0	7.9	1.78	
1	300000000.0	2007.0	5000.0	7.1	2.35	
2	245000000.0	2015.0	393.0	6.8	2.35	
3	250000000.0	2012.0	23000.0	8.5	2.35	
4	0.0	0.0	12.0	7.1	0.00	
...	
5039	0.0	0.0	593.0	7.5	16.00	
5040	1400.0	2013.0	0.0	6.3	0.00	
5041	0.0	2012.0	719.0	6.3	2.35	
5042	1100.0	2004.0	23.0	6.6	1.85	
5043	0.0	2014.0	0.0	0.0	0.00	

	movie_facebook_likes	sales
0	33000.0	523505847.0
1	0.0	9404152.0
2	85000.0	-44925825.0

3	164000.0	198130642.0
4	0.0	0.0
...
5039	32000.0	0.0
5040	16.0	-1400.0
5041	660.0	10443.0
5042	456.0	84122.0
5043	0.0	0.0

[5044 rows x 30 columns]

0.4 Which directors garnered the most total sales?

```
[17]: # gets sum sales of each director and retrieves the director with the most
      ↪total sales
directors_most_sales = data.groupby("director_name")["sales"].sum().
      ↪sort_values(ascending = False)
directors_most_sales[0:1]
```

```
[17]: director_name
Steven Spielberg    2.451332e+09
Name: sales, dtype: float64
```

0.5 Answer: Steven Spielberg


0.6 Plot sales and average likes as a scatterplot. Fit it with a line.

```
[67]: data['avg_likes'] = data[['director_facebook_likes', 'actor_1_facebook_likes',
      ↪'actor_2_facebook_likes', 'actor_3_facebook_likes',
      ↪'movie_facebook_likes', 'cast_total_facebook_likes']].mean(axis=1)

data.duration.quantile(0.99)
data['sales'] = np.clip(data['sales'], 0, 197243982.28999966)
# creates a scatterplot using sales and average_likes
fig = plt.figure(figsize=(15,8))
plt.scatter(data["sales"], data["avg_likes"])

# creates line of best fit for scatterplot
m, b = np.polyfit(data["sales"], data["avg_likes"], 1)
plt.plot(data["sales"], m*data["sales"] + b, color='red')

plt.show()
data.sales.quantile(0.99)
```



[67]: 196952019.9052996

0.7 Which of these genres are the most profitable? Plot their sales using different histograms, superimposed in the same axis.


- Romance
- Comedy
- Action
- Fantasy

```
[21]: # Filters data based on whether it contains the selected genre
romance = data['genres'].str.contains('Romance')
comedy = data['genres'].str.contains('Comedy')
action = data['genres'].str.contains('Action')
fantasy = data['genres'].str.contains('Fantasy')

# Plots histogram from genres
ax = sns.histplot(data["sales"][romance],bins=20, color="red",
    ↪label="Romance",linewidth=0, alpha=0.2, )
sns.histplot(data["sales"][comedy],bins=20, color="teal", label="Comedy",
    ↪linewidth=0, ax=ax, alpha=0.2)
sns.histplot(data["sales"][action],bins=20, color="blue", label="Action",
    ↪linewidth=0, ax=ax, alpha=0.2)
sns.histplot(data["sales"][fantasy],bins=20, color="green", label="Fantasy",
    ↪linewidth=0, ax=ax, alpha=0.2)

ax.legend()
```

[21]: <matplotlib.legend.Legend at 0x1c281616e10>



0.8 For each of movie, compute average likes of the three actors and store it as a new variable

Read up on the mean function.

Store it as a new column, average_actor_likes.

```
[23]: data['average_actor_likes'] = data[['actor_1_facebook_likes',  
    ↪ 'actor_2_facebook_likes', 'actor_3_facebook_likes']].mean(axis=1)  
data
```

```
[23]:
```

	Unnamed: 0	movie_title	color \
0	0	b'Avatar'	Color
1	1	b"Pirates of the Caribbean: At World's End"	Color
2	2	b'Spectre'	Color
3	3	b'The Dark Knight Rises'	Color
4	4	b'Star Wars: Episode VII - The Force Awakens ...	0
...
5039	5039	b'The Following	Color
5040	5040	b'A Plague So Pleasant'	Color

5041	5041	b'Shanghai Calling'	Color
5042	5042	b'My Date with Drew'	Color
5043	5043	b'Starting Over Again'	0

	director_name	num_critic_for_reviews	duration	\
0	James Cameron	723.0	178.0	
1	Gore Verbinski	302.0	169.0	
2	Sam Mendes	602.0	148.0	
3	Christopher Nolan	813.0	164.0	
4	Doug Walker	0.0	0.0	
...	
5039	0	43.0	43.0	
5040	Benjamin Roberds	13.0	76.0	
5041	Daniel Hsia	14.0	100.0	
5042	Jon Gunn	43.0	90.0	
5043	Olivia Lamasan	0.0	0.0	

	director_facebook_likes	actor_3_facebook_likes	actor_2_name	\
0	0.0	855.0	Joel David Moore	
1	563.0	1000.0	Orlando Bloom	
2	0.0	161.0	Rory Kinnear	
3	22000.0	23000.0	Christian Bale	
4	131.0	0.0	Rob Walker	
...	
5039	0.0	319.0	Valorie Curry	
5040	0.0	0.0	Maxwell Moody	
5041	0.0	489.0	Daniel Henney	
5042	16.0	16.0	Brian Herzlinger	
5043	0.0	0.0	Toni Gonzaga	

	actor_1_facebook_likes	...	content_rating	budget	title_year	\
0	1000.0	...	PG-13	237000000.0	2009.0	
1	40000.0	...	PG-13	300000000.0	2007.0	
2	11000.0	...	PG-13	245000000.0	2015.0	
3	27000.0	...	PG-13	250000000.0	2012.0	
4	131.0	...	0	0.0	0.0	
...	
5039	841.0	...	TV-14	0.0	0.0	
5040	0.0	...	0	1400.0	2013.0	
5041	946.0	...	PG-13	0.0	2012.0	
5042	86.0	...	PG	1100.0	2004.0	
5043	0.0	...	PG	0.0	2014.0	

	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes	\
0	936.0	7.9	1.78	33000.0	
1	5000.0	7.1	2.35	0.0	
2	393.0	6.8	2.35	85000.0	

3	23000.0	8.5	2.35	164000.0
4	12.0	7.1	0.00	0.0
...
5039	593.0	7.5	16.00	32000.0
5040	0.0	6.3	0.00	16.0
5041	719.0	6.3	2.35	660.0
5042	23.0	6.6	1.85	456.0
5043	0.0	0.0	0.00	0.0

	sales	avg_likes	average_actor_likes
0	1.972440e+08	6770.833333	930.333333
1	9.404152e+06	15818.833333	15333.333333
2	0.000000e+00	18042.333333	3851.333333
3	1.972440e+08	60959.833333	24333.333333
4	0.000000e+00	69.500000	47.666667
...
5039	0.000000e+00	5917.666667	584.333333
5040	0.000000e+00	2.666667	0.000000
5041	1.044300e+04	866.666667	718.000000
5042	8.412200e+04	126.666667	41.666667
5043	0.000000e+00	0.000000	0.000000

[5044 rows x 32 columns]

0.9 Copying the whole dataframe

```
[25]: df = data.copy()
df.head()
```

```
[25]: Unnamed: 0      movie_title  color \
0      0      b'Avatar'  Color
1      1  b"Pirates of the Caribbean: At World's End"  Color
2      2      b'Spectre'  Color
3      3      b'The Dark Knight Rises'  Color
4      4  b'Star Wars: Episode VII - The Force Awakens ...  0
```

	director_name	num_critic_for_reviews	duration	\
0	James Cameron	723.0	178.0	
1	Gore Verbinski	302.0	169.0	
2	Sam Mendes	602.0	148.0	
3	Christopher Nolan	813.0	164.0	
4	Doug Walker	0.0	0.0	

	director_facebook_likes	actor_3_facebook_likes	actor_2_name	\
0	0.0	855.0	Joel David Moore	
1	563.0	1000.0	Orlando Bloom	
2	0.0	161.0	Rory Kinnear	

3	22000.0	23000.0	Christian Bale
4	131.0	0.0	Rob Walker

	actor_1_facebook_likes	...	content_rating	budget	title_year	\
0	1000.0	...	PG-13	237000000.0	2009.0	
1	40000.0	...	PG-13	300000000.0	2007.0	
2	11000.0	...	PG-13	245000000.0	2015.0	
3	27000.0	...	PG-13	250000000.0	2012.0	
4	131.0	...	0	0.0	0.0	

	actor_2_facebook_likes	imdb_score	aspect_ratio	movie_facebook_likes	\
0	936.0	7.9	1.78	33000.0	
1	5000.0	7.1	2.35	0.0	
2	393.0	6.8	2.35	85000.0	
3	23000.0	8.5	2.35	164000.0	
4	12.0	7.1	0.00	0.0	

	sales	avg_likes	average_actor_likes
0	1.972440e+08	6770.833333	930.333333
1	9.404152e+06	15818.833333	15333.333333
2	0.000000e+00	18042.333333	3851.333333
3	1.972440e+08	60959.833333	24333.333333
4	0.000000e+00	69.500000	47.666667

[5 rows x 32 columns]

0.10 Min-Max Normalization

Normalization is a technique often applied as part of data preparation for machine learning. The goal of normalization is to change the values of numeric columns in the dataset to a common scale, without distorting differences in the ranges of values. For machine learning, every dataset does not require normalization. It is required only when features have different ranges.

The min-max approach (often called normalization) rescales the feature to a hard and fast range of [0,1] by subtracting the minimum value of the feature then dividing by the range. We can apply the min-max scaling in Pandas using the `.min()` and `.max()` methods.

$$x_{scaled} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

0.10.1 Normalize each numeric column (those that have types integer or float) of the copied dataframe (df)

```
[39]: df.dtypes
```

```
[39]: Unnamed: 0          int64
      movie_title      object
      color          object
```

```

director_name          object
num_critic_for_reviews float64
duration               float64
director_facebook_likes float64
actor_3_facebook_likes float64
actor_2_name           object
actor_1_facebook_likes float64
gross                 float64
genres                 object
actor_1_name           object
num_voted_users        float64
cast_total_facebook_likes float64
actor_3_name           object
facenumber_in_poster   float64
plot_keywords          object
movie_imdb_link         object
num_user_for_reviews   float64
language               object
country                object
content_rating          object
budget                 float64
title_year             float64
actor_2_facebook_likes float64
imdb_score             float64
aspect_ratio           float64
movie_facebook_likes    float64
sales                  float64
average_actor_likes     float64
dtype: object

```

```

[27]: max = df['num_critic_for_reviews'].max()
      min = df['num_critic_for_reviews'].min()
      df['nm_num_critic_for_reviews'] = (df['num_critic_for_reviews'] - min) / (
      ↪(max-min))

```

```

[29]: max = df['duration'].max()
      min = df['duration'].min()
      df['nm_duration'] = (df['duration'] - min) / (max-min)

```

```

[31]: max = df['director_facebook_likes'].max()
      min = df['director_facebook_likes'].min()
      df['nm_director_facebook_likes'] = (df['director_facebook_likes'] - min) / (
      ↪(max-min))

```

```

[33]: max = df['actor_3_facebook_likes'].max()
      min = df['actor_3_facebook_likes'].min()

```

```
df['nm_actor_3_facebook_likes'] = (df['actor_3_facebook_likes'] - min) /  $\sqrt{\text{max-min}}$ 
```

```
[35]: max = df['actor_1_facebook_likes'].max()  
min = df['actor_1_facebook_likes'].min()  
df['nm_actor_1_facebook_likes'] = (df['actor_1_facebook_likes'] - min) /  $\sqrt{\text{max-min}}$ 
```

```
[37]: max = df['gross'].max()  
min = df['gross'].min()  
df['nm_gross'] = (df['gross'] - min) / (max-min)
```

```
[39]: max = df['num_voted_users'].max()  
min = df['num_voted_users'].min()  
df['nm_num_voted_users'] = (df['num_voted_users'] - min) / (max-min)
```

```
[41]: max = df['cast_total_facebook_likes'].max()  
min = df['cast_total_facebook_likes'].min()  
df['nm_cast_total_facebook_likes'] = (df['cast_total_facebook_likes'] - min) /  $\sqrt{\text{max-min}}$ 
```

```
[43]: max = df['facenumber_in_poster'].max()  
min = df['facenumber_in_poster'].min()  
df['nm_facenumber_in_poster'] = (df['facenumber_in_poster'] - min) / (max-min)
```

```
[45]: max = df['num_user_for_reviews'].max()  
min = df['num_user_for_reviews'].min()  
df['nm_num_user_for_reviews'] = (df['num_user_for_reviews'] - min) / (max-min)
```

```
[47]: max = df['budget'].max()  
min = df['budget'].min()  
df['nm_budget'] = (df['budget'] - min) / (max-min)
```

```
[49]: max = df['title_year'].max()  
min = df['title_year'].min()  
df['nm_title_year'] = (df['title_year'] - min) / (max-min)
```

```
[51]: max = df['actor_2_facebook_likes'].max()  
min = df['actor_2_facebook_likes'].min()  
df['nm_actor_2_facebook_likes'] = (df['actor_2_facebook_likes'] - min) /  $\sqrt{\text{max-min}}$ 
```

```
[53]: max = df['imdb_score'].max()  
min = df['imdb_score'].min()  
df['nm_imdb_score'] = (df['imdb_score'] - min) / (max-min)
```

```
[55]: max = df['aspect_ratio'].max()
min = df['aspect_ratio'].min()
df['nm_aspect_ratio'] = (df['aspect_ratio'] - min) / (max-min)
```

```
[57]: max = df['movie_facebook_likes'].max()
min = df['movie_facebook_likes'].min()
df['nm_movie_facebook_likes'] = (df['movie_facebook_likes'] - min) / (max-min)
```

```
[59]: max = df['sales'].max()
min = df['sales'].min()
df['nm_sales'] = (df['sales'] - min) / (max-min)
df
```

```
[59]: Unnamed: 0      movie_title  color \
0      0      b'Avatar'  Color
1      1  b"Pirates of the Caribbean: At World's End"  Color
2      2      b'Spectre'  Color
3      3  b'The Dark Knight Rises'  Color
4      4  b'Star Wars: Episode VII - The Force Awakens ...    0
...      ...      ...      ...
5039    5039      b'The Following'  Color
5040    5040  b'A Plague So Pleasant'  Color
5041    5041  b'Shanghai Calling'  Color
5042    5042  b'My Date with Drew'  Color
5043    5043  b'Starting Over Again'    0
```

```
      director_name  num_critic_for_reviews  duration \
0      James Cameron      723.0      178.0
1      Gore Verbinski      302.0      169.0
2      Sam Mendes      602.0      148.0
3  Christopher Nolan      813.0      164.0
4      Doug Walker      0.0      0.0
...      ...      ...      ...
5039      0      43.0      43.0
5040  Benjamin Roberds      13.0      76.0
5041      Daniel Hsia      14.0      100.0
5042      Jon Gunn      43.0      90.0
5043  Olivia Lamasan      0.0      0.0
```

```
      director_facebook_likes  actor_3_facebook_likes      actor_2_name \
0      0.0      855.0  Joel David Moore
1      563.0      1000.0  Orlando Bloom
2      0.0      161.0  Rory Kinnear
3      22000.0      23000.0  Christian Bale
4      131.0      0.0  Rob Walker
...      ...      ...      ...
5039      0.0      319.0  Valorie Curry
```

5040	0.0	0.0	Maxwell Moody
5041	0.0	489.0	Daniel Henney
5042	16.0	16.0	Brian Herzlinger
5043	0.0	0.0	Toni Gonzaga

	actor_1_facebook_likes	nm_cast_total_facebook_likes	\
0	1000.0	...	0.007361
1	40000.0	...	0.073622
2	11000.0	...	0.017816
3	27000.0	...	0.162561
4	131.0	...	0.000218
...
5039	841.0	...	0.002669
5040	0.0	...	0.000000
5041	946.0	...	0.003633
5042	86.0	...	0.000248
5043	0.0	...	0.000000

	nm_facenumber_in_poster	nm_num_user_for_reviews	nm_budget	\
0	0.000000	0.603557	1.940158e-02	
1	0.000000	0.244664	2.455896e-02	
2	0.023256	0.196443	2.005649e-02	
3	0.000000	0.533794	2.046580e-02	
4	0.000000	0.000000	0.000000e+00	
...	
5039	0.023256	0.070949	0.000000e+00	
5040	0.000000	0.000593	1.146085e-07	
5041	0.116279	0.001779	0.000000e+00	
5042	0.000000	0.016601	9.004953e-08	
5043	0.000000	0.000000	0.000000e+00	

	nm_title_year	nm_actor_2_facebook_likes	nm_imdb_score	nm_aspect_ratio	\
0	0.996528	0.006832	0.831579	0.111250	
1	0.995536	0.036496	0.747368	0.146875	
2	0.999504	0.002869	0.715789	0.146875	
3	0.998016	0.167883	0.894737	0.146875	
4	0.000000	0.000088	0.747368	0.000000	
...	
5039	0.000000	0.004328	0.789474	1.000000	
5040	0.998512	0.000000	0.663158	0.000000	
5041	0.998016	0.005248	0.663158	0.146875	
5042	0.994048	0.000168	0.694737	0.115625	
5043	0.999008	0.000000	0.000000	0.000000	

	nm_movie_facebook_likes	nm_sales
0	0.094556	1.000000
1	0.000000	0.047678

2	0.243553	0.000000
3	0.469914	1.000000
4	0.000000	0.000000
...
5039	0.091691	0.000000
5040	0.000046	0.000000
5041	0.001891	0.000053
5042	0.001307	0.000426
5043	0.000000	0.000000

[5044 rows x 49 columns]